Custom Representations of Inductive Families

Constantine Theocharis[0000−1111−2222−3333] and
Edwin Brady[0000−1111−2222−3333]

University of St Andrews, UK
{kt81,ecb10}@st-andrews.ac.uk

Abstract. Inductive families enable a convenient way of programming with dependent types, where dependent pattern matching automatically handles a lot of the otherwise manual work of reasoning about index refinement. However, when it comes to compilation, the default representation of inductive types can be inefficient. Often, we want multiple indexed views of the same data, which has runtime overhead with current methods. In this paper, introduce a language with dependent types, and inductive families with custom representations. A type family F can be used to represent the inductive family D as long as T can mimic D's constructors and induction principle. This way, we can build up a standard library of convenient inductive families, whose re-indexing conversion functions are erased at compile-time. In particular, we show how we can reproduce and extend the 'Nat-hack' from Agda, Idris and Lean, fully within the language. Since we are working with dependent types, we can indeed reason about data representations internally; in this spirit, we are awarded with a correctness proof of the Nat-hack, for free.

1 Introduction

Inductive families are a broad generalisation of inductive data types found in most functional programming languages. Every inductive definition is equipped with an induction principle that captures the notion of mathematical induction over the data, and in particular, enables structural recursion over the data. This is a powerful tool for programming as well as theorem proving. However, this abstraction often comes at a cost; the data representation of inductive types is a linked tree structure. By default, a list as an inductively defined data structure is stored as a linked list, and a natural number is stored as a unary number where each digit is an empty heap cell. This representation is not always the most efficient for all operations, and often forces users to rely on 'more efficient' machine primitives to achieve good performance. The 'Nat-hack' in languages with dependent types is a prime example of this, where natural numbers are represented as big integers for efficient computation.

In this paper, we propose an extension to a core language with dependent types and inductive families which allows users to define custom representations for inductive types. This is done through a translation of the constructors and eliminators of the inductive type to a concrete implementation, which forms

a 'representation'. Representations are defined in the language itself, and come with coherence properties that ensure that the representation is a strong bijection between the inductive type and its implementation.

2   A tour of data representations

2.1   Natural numbers

A common optimisation done by proof assistants such as Idris2 and Lean is to represent natural numbers as GMP-style big integers. By virtue of their definition, natural numbers are represented as unary numbers

$$\textbf{data } \mathsf{Nat} \left\{ \begin{array}{l} 0 : \mathsf{Nat} \\ 1+ : \mathsf{Nat} \to \mathsf{Nat} \end{array} \right\}$$

This representation

$$\textbf{record } \mathsf{Player} \left\{ \begin{array}{l} \mathsf{name} : \mathsf{String} \\ \mathsf{health} : \mathsf{Fin\ MAX\text{-}HEALTH} \end{array} \right\}$$

2.2   Views on lists

2.3   Reindexing and forgetful maps for free

Let us assume that in our language, for some data type

$$\textbf{data } \mathsf{D}\ \Delta : \ \varXi \to \mathsf{Type}$$

we have access to an internal description $\mathsf{desc_D} : \Delta \to \mathsf{Desc}\ \varXi$, along with an interpretation function

$$[\![\_]\!] : \mathsf{Desc}\ \varXi \to (\varXi \to \mathsf{Type}) \to (\varXi \to \mathsf{Type})$$

and a fixpoint operator

$$\textbf{data } \mu\,(D : \mathsf{Desc}\ \varXi)\,(\xi : \varXi)\ \left\{ \mathsf{fix} : [\![D]\!]\,(\mu\ D)\ \xi \to \mu\ D\ \xi \right\}$$

such that $\mu\,(\mathsf{desc_D}\ \delta) \simeq \mathsf{D}\ \delta$ is a strong bijection. Then, we might choose to define the representation of D as $\mu\,(\mathsf{desc_D}\ \delta)$; in this way, we represent the 'primitive' datatype D as the interpretation of its code $\mathsf{desc_D}$.

Now consider that we have an indexed datatype

$$\textbf{data } \mathsf{C}\ \Delta : \ (\xi : \varXi) \to \varPhi[\xi] \to \mathsf{Type}$$

which is a refined version of D. If we can construct an algebra

$$\mathsf{phi}\ \delta : [\![\mathsf{desc_D}]\!]\ \varPhi\ \xi \to \varPhi\ \xi$$

that computes the index Φ, then we can form the description

$$\mathsf{alg\text{-}orn}\ (\mathsf{phi}\ \delta) : \mathsf{Desc}\ (\Sigma\ \Xi\ \Phi)$$

which is the ornament induced by the algebra phi. If a strong bijection can be established between $\mu\ (\mathsf{alg\text{-}orn}\ (\mathsf{phi}\ \delta))$ and $\mathsf{C}\ \delta$, then the former can be used as the representation of the latter. Finally, this allows us to define a zero-cost forgetful conversion function from C and D, which is erased at compile-time, as

$$\mathsf{forget\text{-}}\Phi : \mathsf{C}\ \delta\ \xi\ \phi \to \mathsf{D}\ \delta\ \xi \tag{1}$$

$$\mathsf{forget\text{-}}\Phi\ c\ d = \mathsf{unrepr}\ (\mathsf{alg\text{-}orn\text{-}forget}\ (\mathsf{repr}\ c)) \tag{2}$$

Example 1. Let $\Delta = (T : \mathsf{Type})$, $\Xi = \cdot$ and $\Phi = (n : \mathsf{Nat})$. Let $\mathsf{D}\ T = \mathsf{List}\ T$, and $\mathsf{C}\ T\ n = \mathsf{Vec}\ T\ n$. Then let $\mathsf{desc_{List}}$ be the description for lists

$$\sigma\ [\mathsf{nil}, \mathsf{cons}]\ \left\{ \begin{array}{l} \mathsf{nil} \mapsto \mathsf{end} \\ \mathsf{cons} \mapsto \sigma\ T\ (\lambda\_.\ \mathsf{node}\times\ \mathsf{end}) \end{array} \right\}$$

and construct the algebra phi from the length function $\mathsf{length} : \mathsf{List}\ T \to \mathsf{Nat}$. It is easy to construct a strong bijection $\mu\ (\mathsf{alg\text{-}orn}\ (\mathsf{length}\ T))\ n \simeq \mathsf{Vec}\ T\ n$, so we can represent the latter as the former, and define a zero-cost function to forget the length of a vector using (1).

2.4   Binary data

3   A type system for data representations

In this section, we describe a type system for data representations in a language with dependent types. We start by defining a core language with dependent types and inductive constructions λ_{IND}. We then extend this language with data representations to form λ_{REP}, which allow users to define custom representations for inductive types and other global symbols. We present these languages with intrinsically well-formed contexts, types, and terms, quotiented by their definitional equality rules [altenkirch].

3.1   A core language with inductive types, λ_{IND}

The core langyge we start with is λ_{IND}. It contains Π-types and a single universe Type with $\mathsf{Type} : \mathsf{Type}$. We are not concerned with universe polymorphism or a sound logical interpretation as this is orthogonal to the main focus of this work. Nevertheless, all the results should be readily extensible to a sound language with a universe hierarchy. We follow a similar approach to [1] by packaging named inductive constructions and global function definitions into a signature, and indexing contexts by signatures. A typing judgement looks like

$$\Sigma \mid \Gamma \vdash t : T$$

$$\frac{}{\cdot \ \mathrm{sig}} \ \text{Sig-Empty}$$

$$\text{Sig-Extend} \quad \frac{\Sigma \ \mathrm{sig} \qquad \Sigma \vdash Z}{\Sigma, Z \ \mathrm{sig}}$$

$$\text{Con-Empty} \quad \frac{\Sigma \ \mathrm{sig}}{\Sigma \vdash \cdot \ \mathrm{con}}$$

$$\text{Con-Extend} \quad \frac{\Sigma \vdash \Gamma \ \mathrm{con} \qquad \Sigma \mid \Gamma \vdash T \ \mathrm{type}}{\Sigma \vdash \Gamma, T \ \mathrm{con}}$$

$$\text{Tel-Empty} \quad \frac{\Sigma \vdash \Gamma \ \mathrm{con}}{\Sigma \mid \Gamma \vdash \cdot \ \mathrm{tel}}$$

$$\text{Tel-Extend} \quad \frac{\Sigma \mid \Gamma \vdash \Delta \ \mathrm{tel} \qquad \Sigma \mid \Gamma \vdash T \ \mathrm{type}}{\Sigma \mid \Gamma \vdash \Gamma, T \ \mathrm{tel}}$$

Fig. 1. Rules for signatures, contexts and telescopes in λ_{IND}.

and is read as "in signature Σ and context Γ, term t has type T". The rules for signatures, contexts and telescopes are given in fig. 1.

Telescopes [deBrujin] are very similar to contexts, but restricted to types from a single stage and well formed with respect to a context Γ, meaning that telescopes can contain open terms. We use the notation $\Delta \to t$ to denote a repeated function type with parameters from Δ and codomain T which may depend on the parameters. Additionally, we will sometimes explicitly bind the names of a telescope such as $(\delta : \Delta) \to T[\delta]$. Similar syntax is used to extend contexts with telescopes: Γ, Δ or $\Gamma, \delta : \Delta$.

Awkward spacing!

Next, the rules for well-formed items in signatures are given in fig. 2.

$$\text{Data-Item} \quad \frac{\Sigma \mid \cdot \vdash \Delta \ \mathrm{tel} \qquad \Sigma \mid \Delta \vdash \Xi \ \mathrm{tel} \qquad \mathsf{D} \ \mathrm{label} \notin \Sigma}{\Sigma \vdash \mathbf{data} \ \mathsf{D} \ \Delta : \Xi \to \mathsf{Type}}$$

$$\text{Ctor-Item} \quad \frac{\mathbf{data} \ \mathsf{D} \ \Delta : \Xi \to \mathsf{Type} \in \Sigma}{\mathbf{closed} \ \mathsf{D}, \_ \notin \Sigma \quad \mathsf{C} \ \mathrm{label} \notin \Sigma \quad \Sigma \mid \Delta \vdash \Pi \ \mathrm{tel} \quad \Sigma \mid \Delta, \Pi \vdash \xi : \Xi}{\Sigma \vdash \mathbf{ctor} \ \mathsf{C} \ \Pi : \mathsf{D} \ \Delta \ \xi}$$

$$\text{Closed-Item} \quad \frac{\mathbf{data} \ \mathsf{D} \ \Delta : \Xi \to \mathsf{Type} \in \Sigma \quad \forall i \in I. \ \mathbf{ctor} \ \mathsf{C}_i \ \Pi_i : \mathsf{D} \ \Delta \ \xi_i \in \Sigma \quad \mathbf{closed} \ \mathsf{D}, \_ \notin \Sigma}{\Sigma \vdash \mathbf{closed} \ \mathsf{D}, \mathsf{C}}$$

$$\text{Def-Item} \quad \frac{\Sigma \mid \cdot \vdash m : M \qquad \mathsf{f} \ \mathrm{label} \notin \Sigma}{\Sigma \vdash \mathbf{def} \ \mathsf{f} : M = m}$$

Fig. 2. Rules for items in signatures in λ_{IND}.

DATA-FORM
$$\frac{\textbf{data } \mathsf{D} \, \Delta : \Xi \to \mathsf{Type} \in \Sigma}{\Sigma \mid \Gamma \vdash \mathsf{D} : \Delta \to \Xi \to \mathsf{Type}}$$

DATA-INTRO
$$\frac{\textbf{data } \mathsf{D} \, \Delta : \Xi \to \mathsf{Type} \in \Sigma \qquad \textbf{ctor } \mathsf{C} \, \Pi : \mathsf{D} \, \Delta \, \xi \in \Sigma}{\Sigma \mid \Gamma \vdash \mathsf{C} : (\delta : \Delta) \to (\pi : \Pi[\delta]) \to \mathsf{D} \, \delta \, (\xi[\delta, \pi])}$$

DATA-CASE
$$\frac{\textbf{data } \mathsf{D} \, \Delta : \Xi \to \mathsf{Type} \in \Sigma \qquad \forall i \in I. \ \textbf{ctor } \mathsf{C}_i \, \Pi_i : \mathsf{D} \, \xi_i \in \Sigma \qquad \textbf{closed } \mathsf{D} \, C \in \Sigma}{\Sigma \mid \Gamma \vdash \mathsf{case}_{\mathsf{D}} : (\delta : \Delta) \to \mathrm{Case}((\xi : \Xi[\delta], x : \mathsf{D} \, \delta \, \xi), \{(\pi : \Pi_i[\delta]), (\xi_i[\delta], \mathsf{C}_i \, \delta \, \pi)\}_i)}$$

DATA-CASE-ID$_j$
$$\frac{\textbf{data } \mathsf{D} \, \Delta : \Xi \to \mathsf{Type} \in \Sigma \qquad \qquad \qquad \\ \forall i \in I. \ \textbf{ctor } \mathsf{C}_i \, \Pi_i : \mathsf{D} \, \xi_i \in \Sigma \qquad \textbf{closed } \mathsf{D} \, C \in \Sigma \qquad \Sigma \mid \Gamma \vdash \delta : \Delta}{\Sigma \mid \Gamma \vdash \mathrm{ValidCase}(\mathsf{case}_{\mathsf{D}} \, \delta)}$$

DEF-INTRO
$$\frac{\textbf{def } \mathsf{f} : M = m \in \Sigma}{\Sigma \mid \Gamma \vdash \mathsf{f} : M}$$

Fig. 3. Terms and types associated to items in signatures in λ_{IND}.

The rules for data constructors do not consider the recursive occurences of D explicitly, which means that strict positivity is not ensured. Rather, we assume that a separate check is performed to ensure that the defined types adhere to strict positivity, if necessary. As a result, the eliminator for data types does not provide the inductive hypotheses directly, but we assume that the language allows general recursion. Similarly to the positivity requirements, we expect that if termination is a desirable property of the system, it is ensured separately to the provided typing rules.

In fig. 3, the type $\Sigma \mid \Gamma \vdash \mathrm{Case}(\Phi, \{\Pi_k, \phi_k\}_k) : \mathsf{Type}$ is defined as

$$\frac{\Sigma \mid \Gamma \vdash \Phi \ \mathrm{tel} \qquad \forall k \in K. \ \Sigma \mid \Gamma \vdash \Pi_k \ \mathrm{tel} \qquad \forall k \in K. \ \Sigma \mid \Gamma, \Pi_k \vdash \phi_k : \Phi}{\mathrm{Case}(\Phi, \{\Pi_k, \phi_k\}_k) := (P : \Phi \to \mathsf{Type}) \to \{(\Pi_k \to P \, \phi_k)\}_k \to (\phi : \Phi) \to P\phi}$$

The $\Sigma \mid \Gamma \vdash \mathrm{ValidCase}(c)$ condition is defined inductively as

$$\frac{\Sigma \mid \Gamma \vdash c : \mathrm{Case}(\Phi, \{\Pi_k, \phi_k\}_k) \\ \forall j \in K. \ \Sigma \mid \Gamma, P : \Phi \to \mathsf{Type}, \{\kappa_k : \Pi_k \to P \, \phi_k\}_k, \pi : \Pi_j \\ \vdash c \, P \, \{\kappa_k\}_k \, \phi_j[\pi] \equiv \kappa_j \, \pi : P \, \phi_j[\pi]}{\Sigma \mid \Gamma \vdash \mathrm{ValidCase}(c)}$$

3.2 Extending λ_{IND} with data representations

We base language λ_{IND} to form λ_{REP}, which allows users to define custom representations for inductive types and global functions.

$$\text{Repr-Data}$$
$$\frac{\textbf{data } \mathsf{D}\,\Delta : \Xi \to \mathsf{Type} \in \Sigma \qquad \Sigma \mid \Delta, \Xi \vdash A : \mathsf{Type}}{\Sigma \vdash \textbf{repr } \mathsf{D}\,\Delta\,\Xi \textbf{ as } A}$$

$$\text{Repr-Ctor}$$
$$\frac{\textbf{data } \mathsf{D}\,\Delta : \Xi \to \mathsf{Type} \in \Sigma \qquad \forall i \in I.\ \textbf{ctor } \mathsf{C}_i\,\Pi_i : \mathsf{D}\,\Delta\,\xi_i \in \Sigma \qquad \textbf{repr } \mathsf{D}\,\Delta\,\Xi \textbf{ as } A \in \Sigma \qquad \forall i < k.\ \textbf{repr } \mathsf{C}_i\,\Pi_i \textbf{ as } t_i \in \Sigma \qquad \Sigma \mid \Delta, \Pi_k \vdash t_k : A\,\Delta\,\xi_k}{\Sigma \vdash \textbf{repr } \mathsf{C}_k\,\Pi_k \textbf{ as } t_k}$$

$$\text{Repr-Case}$$
$$\frac{\begin{array}{c}\textbf{data } \mathsf{D}\,\Delta : \Xi \to \mathsf{Type} \in \Sigma \\ \textbf{repr } \mathsf{D}\,\Delta\,\Xi \textbf{ as } A \in \Sigma \qquad \forall i \in I.\ \textbf{ctor } \mathsf{C}_i\,\Pi_i : \mathsf{D}\,\Delta\,\xi_i \in \Sigma \\ \forall i \in I.\ \textbf{repr } \mathsf{C}_i\,\Pi_i \textbf{ as } t_i \in \Sigma \qquad \textbf{closed } \mathsf{D}\,\mathsf{C} \in \Sigma \\ \Sigma \mid \delta : \Delta \vdash c : \mathrm{Case}((\xi : \Xi[\delta], x : A\,\delta\,\xi), \{(\pi : \Pi_i[\delta]), (\xi_i[\delta], A\,\delta\,\pi)\}_i) \\ \Sigma \mid \Delta \vdash \mathrm{ValidCase}(c)\end{array}}{\Sigma \vdash \textbf{repr } \mathsf{case}_\mathsf{D}\,\Delta \textbf{ as } c}$$

$$\text{Repr-Def}$$
$$\frac{\textbf{def } \mathsf{f} : M = m \in \Sigma \qquad \Sigma \mid \cdot \vdash a : M \qquad \Sigma \mid \cdot \vdash m \equiv a : M}{\Sigma \vdash \textbf{repr } \mathsf{f} \textbf{ as } a}$$

Fig. 4. Rules for data representations in λ_{REP}.

3.3 Properties of λ_{REP}

- Decidability of equality
- Confluence of reduction
- Strong normalization

4 Translating from λ_{REP} to λ_{IND}

- Setoid homomorphism
- Faithful translation (setoid injectivity)
- Corollaries

References

1. Cockx, J., Abel, A.: Elaborating dependent (co)pattern matching. Proc. ACM Program. Lang. **2**(ICFP), 1–30 (Jul 2018), https://doi.org/10.1145/3236770