

Custom Representations of Inductive Families

Constantine Theocharis^[0009–0001–0198–2750] and
Edwin Brady^[0000–0002–9734–367X]

University of St Andrews, UK
`{kt81,ecb10}@st-andrews.ac.uk`

Abstract. Inductive families provide a convenient way of programming with dependent types. Yet, when it comes to compilation, their default linked-tree runtime representations, as well as the need to convert between different indexed views of the same data, can lead to unsatisfactory runtime performance. In this paper, we aim to introduce a language with dependent types, and inductive families with custom representations. Representations are a version of Wadler’s views [27], refined to inductive families like in Epigram [25], but with compilation guarantees: a represented inductive family will not leave any runtime traces behind, without relying on heuristics such as deforestation. This way, we can build a library of convenient inductive families based on a minimal set of primitives, whose re-indexing and conversion functions are erased at compile-time. We show how we can express inductive data optimisation techniques, such as representing `Nat`-like types as GMP-style big integers, without special casing in the compiler. With dependent types, reasoning about data representations is also possible: the representation of a type becomes an internal predicate, and we get computationally irrelevant isomorphisms between the original and represented data.

Keywords: Dependent types · Memory representation · Inductive families

1 Introduction

Inductive families are a generalisation of inductive data types found in some programming languages with dependent types. Every inductive definition is equipped with an eliminator that captures the notion of mathematical induction over the data, and in particular, enables structural recursion over the data. This is a powerful tool for programming as well as theorem proving. However, this abstraction has a cost when it comes to compilation: the runtime representation of inductive types is a linked tree structure. This representation is not always the most efficient for all operations, and often forces users to rely on more efficient machine primitives to achieve desirable performance, at the cost of structural recursion and dependent pattern matching. This is the first problem we aim to address in this paper.

Despite advances in the erasure of irrelevant indices in inductive families [14] and the use of theories with irrelevant fragments [10,26], there is still a

need to convert between different indexed views of the same data. For example, the function to convert from `BinTree T` to `BinTreeOfHeight T n` by forgetting the height index n is *not* erased by any current language with dependent types, unless sized binary trees are defined as a refinement of binary trees with an erased height field (which hinders dependent pattern matching due to the presence of non-structural witnesses), or a Church encoding is used in a Curry-style context [18] (which restricts the flexibility of data representation). This is the second problem we aim to address in this paper.

Wadler’s views [27] provide a way to abstract over inductive interfaces, so that different views of the same data can be defined and converted between seamlessly. In the context of inductive families, views have been used in Epigram [25] that utilise the index refinement machinery of dependent pattern matching to avoid certain proof obligations with eliminator-like constructs. While views exhibit a nice way to transport across a bijection between the original data and the viewed data, they do not utilise this bijection to erase the view from the program.

In this paper, we propose an extension `DATATT` to Martin-Löf type theory with which allows programmers to define inductive types with custom, correct-by-construction data representations. This is done through user-defined translations of the constructors and eliminators of an inductive type to a concrete implementation, which form a bijective view of the original data called a ‘representation’. Representations are defined internally to the language, and require coherence properties that ensure a representation is faithful to its the original inductive family. We contribute the following:

- A formulation of common optimisations such as the ‘Nat-hack’, and similarly for other inductive types, as well as zero-cost data reuse when reindexing, using representations (section 2).
- A dependent type system `DATATT` with data types formulated as inductive algebras for signatures, along with a translation to `MLTT` that replaces all data types with their defined inductive algebras (section 3). We have formalised this in Agda.
- A prototype implementation of this system in `SUPERFLUID`, a programming language with inductive types and dependent pattern matching (section 4).

2 A tour of data representations

A common optimisation done by programming languages with dependent types such as Idris 2, Agda, Rocq and Lean is to represent natural numbers as GMP-style big integers. The definition of natural numbers looks like

$$\mathbf{data\ Nat} \left\{ \begin{array}{l} \mathbf{zero} : \mathbf{Nat} \\ \mathbf{succ} : \mathbf{Nat} \rightarrow \mathbf{Nat} \end{array} \right\} \quad (1)$$

and generates a case analysis principle `caseNat` of type

$$(P : \mathbf{Nat} \rightarrow \mathcal{U}) \rightarrow P \mathbf{zero} \rightarrow ((n : \mathbf{Nat}) \rightarrow P (\mathbf{succ} \ n)) \rightarrow (s : \mathbf{Nat}) \rightarrow P \ s.$$

This is a special case of the induction principle `elimNat`, where the inductive hypotheses are also given in each method, but this can be recovered through recursion. Without further intervention, the `Nat` type is represented in unary form, where each digit becomes an empty heap cell at runtime. This is inefficient for a lot of the basic operations on natural numbers, especially since computers are particularly well-equipped to deal with numbers natively, so many real-world implementations will treat `Nat` specially, swapping the default inductive type representation with one based on GMP integers. This is done by performing the replacements

$$\begin{aligned}
 |\text{zero}| &= 0, \\
 |\text{succ}| &= \lambda x \Rightarrow x + 1, \\
 |\text{case}_{\text{Nat}} P m_{\text{zero}} m_{\text{succ}} s| \\
 &= \text{let } s = |s| \text{ in if } s == 0 \\
 &\quad \text{then } |m_{\text{zero}}| \\
 &\quad \text{else } |m_{\text{succ}}| (s - 1),
 \end{aligned}$$

where $|\cdot|$ denotes a source translation into a compilation target language with appropriate big unsigned integer primitives. This is the ‘Nat-hack’.¹

In addition to the constructors and case analysis, the compiler might also define translations for commonly used definitions which have a more efficient counterpart in the target, such as recursively-defined addition, multiplication, etc. The recursively-defined functions are well-suited to proofs and reasoning, while the unsigned integer primitives are more efficient for computation. This way, the surface program can take advantage of the structural properties of induction, notably dependent pattern matching, while still benefiting from an efficient representation at runtime.

The issue with this approach is that it only works for the data types which the compiler can recognise as special. Particularly in the presence of dependent types, other data types might end up being equivalent to `Nat` or another ‘nicely-representable’ type, but in a non-trivial way that the compiler cannot recognise. It is also hard to know when such optimisations will fire and performance can become very brittle upon refactoring. Hence, one of our goals is to extend this optimisation to work for any data type. To achieve this, our framework requires that representations are fully typed in a way that ensures the behaviour of the representation of a data type matches the behaviour of the data type itself.

¹ Idris 2 will in fact look for any `Nat`-like types and apply this optimisation. A similar optimisation is also done with list-like and boolean-like types because they have a canonical representation in the target runtime, Chez Scheme.

2.1 The well-typed Nat-hack

A representation definition looks like

$$\text{repr Nat as UBig} \left\{ \begin{array}{l} \text{zero as } 0 \\ \text{succ } n \text{ as } 1 + n \\ \text{elim}_{\text{Nat}} \text{ as } \text{ubig-elim} \\ \quad \text{by } \text{ubig-elim-zero-id}, \\ \quad \quad \text{ubig-elim-add-one-id} \end{array} \right\}$$

The inductive type `Nat` is represented as the type `UBig` of big unsigned integers, with translations for the constructors `zero` and `succ`, and the eliminator `elimNat` (now with the inductive hypotheses). Additionally, the eliminator is shown to satisfy the expected computation rules of the `Nat` eliminator, which are postulated as propositional equalities. This representation is valid in a context containing the symbols

$$\begin{aligned} 0, 1 &: \text{UBig} & + &: \text{UBig} \rightarrow \text{UBig} \rightarrow \text{UBig} \\ \text{ubig-elim} &: (P : \text{UBig} \rightarrow \mathcal{U}) \rightarrow P \ 0 \rightarrow ((n : \text{UBig}) \rightarrow \overline{P \ n} \rightarrow P \ (1 + n)) \\ &\rightarrow (s : \text{UBig}) \rightarrow P \ s \end{aligned}$$

and propositional equalities

$$\begin{aligned} \text{ubig-elim-zero-id} &:_{\forall P b r} \text{ubig-elim } P \ b \ r \ 0 = b \\ \text{ubig-elim-add-one-id} &:_{\forall P b r n} \text{ubig-elim } P \ b \ r \ (1 + n) = r \ n \ (\lambda _ . \text{ubig-elim } P \ b \ r \ n). \end{aligned}$$

For the remainder of the paper we will work with eliminators rather than case analysis but the approach can be specialised to the latter if the language has general recursion. Assuming call-by-value semantics, inductive hypotheses are labelled \overline{P} which denotes lazy values, that is, functions `Unit` \rightarrow P .

The compiler knows how to perform pattern matching on `Nat`, and produces invocations of `elimNat` as a result. On the other hand, it does not know how to pattern match on `UBig`. With this representation, we get the best of both worlds: `Nat` is used for pattern matching, but is then replaced with `UBig` during compilation, generating more efficient code. We expect that the underlying implementation of `UBig` indeed satisfies these postulated properties, which is a separate concern from the correctness of the representation itself. However, such postulates are only needed when the representation target is a primitive; the next examples use defined types as targets, so that the coherence of the target eliminator follows from the coherence of other eliminators used in its implementation.

2.2 Vectors are just certain lists

In addition to representing inductive types as primitives, we can use representations to share the underlying data when converting between indexed views of

the same data. For example, we can define a representation of `Vec` in terms of `List`, so that the conversion from one to the other is ‘compiled away’. We can do this by representing the indexed type as a refinement of the unindexed type by an appropriate relation. For the case of `Vec`, we know intuitively that

$$\text{Vec } T \ n \simeq \{l : \text{List } T \mid \text{length } l = n\}$$

which we shorthand as $\text{List}' T \ n := \{l : \text{List } T \mid \text{length } l = n\}$. We can thus choose $\text{List}' T \ n$ as the representation of $\text{Vec } T \ n$.² We are then tasked with providing terms that correspond to the constructors of `Vec` but for `List'`. These can be defined as

$$\begin{aligned} \text{nil} &: \text{List}' T \ \text{zero} & \text{cons} &: T \rightarrow \text{List}' T \ n \rightarrow \text{List}' T \ (\text{succ } n) \\ \text{nil} &= (\text{nil}, \text{refl}) & \text{cons } x \ (xs, p) &= (\text{cons } x \ xs, \text{cong } (\text{succ}) \ p) \end{aligned}$$

Next we need to define the eliminator for `List'`, which should have the form

$$\begin{aligned} \text{elim-List}' &: (E : (n : \text{Nat}) \rightarrow \text{List}' T \ n \rightarrow \text{Type}) \\ &\rightarrow E \ \text{zero} \ \text{nil} \\ &\rightarrow ((x : T) \rightarrow (n : \text{Nat}) \rightarrow (xs : \text{List}' T \ n) \rightarrow \overline{E \ n \ xs} \rightarrow E \ (\text{succ } n) \ (\text{cons } x \ xs)) \\ &\rightarrow (n : \text{Nat}) \rightarrow (v : \text{List}' T \ n) \rightarrow E \ n \ v \end{aligned}$$

Dependent pattern matching does a lot of the heavy lifting by refining the length index and equality proof by matching on the underlying list. However we still need to substitute the lemma $\text{cong } (\text{succ}) \ (\text{succ-inj } p) = p$ in the recursive case.

$$\begin{aligned} \text{elim-List}' \ P \ b \ r \ \text{zero} \ (\text{nil}, \text{refl}) &= b \\ \text{elim-List}' \ P \ b \ r \ (\text{succ } m) \ (\text{cons } x \ xs, e) &= \text{subst } (\lambda p. P \ (\text{succ } m) \ (\text{cons } x \ xs, p)) \\ &\quad (\text{succ-cong-id } e) \ (r \ x \ (xs, \text{succ-inj } e)) \\ &\quad (\lambda _ . \text{elim-List}' \ P \ b \ r \ m \ (xs, \text{succ-inj } e))) \end{aligned}$$

Finally, we need to prove that the eliminator satisfies the expected computation rules propositionally. These are

$$\begin{aligned} \text{elim-List}'\text{-nil-id} &: \text{elim-List}' \ P \ b \ r \ \text{zero} \ (\text{nil}, \text{refl}) = b \\ \text{elim-List}'\text{-cons-id} &: \text{elim-List}' \ P \ b \ r \ (\text{succ } m) \ (\text{cons } x \ xs, \text{cong } \text{succ } p) \\ &= r \ x \ (xs, p) \ (\lambda _ . \text{elim-List}' \ P \ b \ r \ m \ (xs, p)) \end{aligned}$$

The first holds definitionally, and the second requires a small amount of equality reasoning if the theory does not have definitional uniqueness of identity. This completes the definition of the representation of `Vec` as `List'`, which would be

² We will take the subset $\{x : A \mid P \ x\}$ to mean a Σ -type $(x : A) \times P \ x$ where the right component is irrelevant and erased at runtime.

written as

$$\text{repr } \text{Vec } T \ n \text{ as } \text{List}' \ T \ n \left\{ \begin{array}{l} \text{nil as nil} \\ \text{cons as cons} \\ \text{elim}_{\text{Vec}} \text{ as elim-List}' \\ \text{by elim-List'-nil-id,} \\ \text{elim-List'-cons-id} \end{array} \right\}$$

Now the hard work is done; Every time we are working with a $v : \text{Vec } T \ n$, its form will be (l, p) at runtime, where l is the underlying list and p is the proof that the length of l is n . Under the assumption that the Σ -type's right component is irrelevant and erased at runtime, every vector is simply a list at runtime, where the length proof has been erased. In practice, this erasure is achieved in SUPERFLUID using quantitative type theory [10]. In section 3.7 we show how to formally identify computationally irrelevant conversion functions.

We can utilise this representation to convert between **Vec** and **List** at zero runtime cost, by using the **repr** and **unrepr** operators of the language (defined in section 3). Specifically, we can define the functions

```
forget-length : Vec T n → List T
forget-length v = let (l, _) = repr v in l

remember-length : (l : List T) → Vec T (length l)
remember-length l = unrepr (l, refl)
```

and in section 3.7 we will show that it holds by reflexivity that such functions are inverses of one another.

2.3 General reindexing

The idea from the previous example can be generalised to any data type. In general, suppose that we have two inductive families

$$\mathbf{F} : P \rightarrow \mathcal{U} \quad \mathbf{G} : (p : P) \rightarrow X \ p \rightarrow \mathcal{U}$$

for some index family $X : P \rightarrow \mathcal{U}$. If we hope to represent **G** as some refinement of **F** then we must be able to provide a way to compute **G**'s extra indices X from **F**, like we computed **Vec**'s extra **Nat** index from **List** with **length** in the previous example. This means that we need to provide a function **comp** : $\forall p. \mathbf{F} \ p \rightarrow X \ p$ which can then be used to form the family

$$\mathbf{F}^{\text{comp}} \ p \ x := \{f : \mathbf{F} \ p \mid \text{comp } f = x\}.$$

If **G** is 'equivalent' to the algebraic ornament of **F** by the algebra defining **comp** (given by an isomorphism between the underlying polynomial functors), then it is also equivalent to the Σ -type above. The 'recomputation lemma' of algebraic ornaments [17] then arises from its projections. Our system allows us to *set* the representation of **G** as \mathbf{F}^{comp} , so that the forgetful map from **G** to **F** is the identity at runtime.

2.4 Zero-copy deserialisation

The machinery of representations can be used to implement zero-copy deserialisation of data formats into inductive types. Here we sketch how this works. Consider the following record for a player in a game:

$$\text{data Player} \left\{ \begin{array}{l} \text{player} : (\text{position} : \text{Position}) \\ \quad \rightarrow (\text{direction} : \text{Direction}) \\ \quad \rightarrow (\text{items} : \text{Fin MAX_INVENTORY}) \\ \quad \rightarrow (\text{inventory} : \text{Inventory (fin-to-nat items)}) \rightarrow \text{Player} \end{array} \right\}$$

We can use the `Fin` type to maintain the invariant that the inventory has a maximum size. Additionally, we can index the `Inventory` type by the number of items it contains, which might be defined similarly to `Vec`:

$$\text{data Inventory } (n : \text{Nat}) \left\{ \begin{array}{l} \text{empty} : \text{Inventory } 0 \\ \text{add} : \text{Item} \rightarrow \text{Inventory } n \rightarrow \text{Inventory } (1 + n) \end{array} \right\}$$

We can use the full power of inductive families to model the domain of our problem in the way that is most convenient for us. If we were writing this in a lower-level language, we might choose to use the serialised format directly when manipulating the data, relying on the appropriate pointer arithmetic to access the fields of the serialised data, to avoid copying overhead. Representations allow us to do this while still being able to work with the high-level inductive type.

We can define a representation for `Player` as a pair of a byte buffer and a proof that the byte buffer contents correspond to a player record. Similarly, we can define a representation for `Inventory` as a pair of a byte buffer and a proof that the byte buffer contents correspond to an inventory record of a certain size. By the implementation of the eliminator in the representation, the projection `inventory : (p : Player) → Inventory p.items` is compiled into some code to slice into the inventory part of the player's byte buffer. We assume that the standard library already represents `Fin` in the same way as `Nat`, so that reading the `items` field is a constant-time operation (we do not need to build a unary numeral). We can thus define the representation of `Player` as

$$\text{repr Player as } \{\text{Buf} \mid \text{IsPlayer}\} \left\{ \begin{array}{l} \text{player as buf-is-player} \\ \text{elim}_{\text{Player}} \text{ as elim-buf-is-player} \\ \quad \text{by elim-buf-is-player-id} \end{array} \right\}$$

with an appropriate definition of `IsPlayer` which refines a byte buffer. The refinement would have to match the expected structure of the byte buffer, so that all the required fields can be extracted. Allais [8] explores how data descriptions that index into a flat buffer can be defined.

2.5 Transitivity

Representations are transitive, so in the previous example, the eventual representation of `Vec` at runtime is determined by the representation of `List`. It is possible

to define a custom representation for `List` itself, for example a heap-backed array or a finger tree, and `Vec` would inherit this representation. However it will still be the case that $\text{Repr } (\text{Vec } T \ n) \equiv \text{List } T$, which means the `repr`/`Repr` operators only look at the immediate defined representation of a term. Regardless, we can construct predicates that find types which satisfy a certain eventual representation. For example, given a `Buf` type of byte buffers, we can consider the set of all types which are eventually represented as a `Buf`:

$$\text{data ReprBuf } (T : \mathcal{U}) \left\{ \begin{array}{l} \text{buf} : \text{ReprBuf Buf} \\ \text{from} : \text{ReprBuf } (\text{Repr } T) \rightarrow \text{ReprBuf } T \\ \text{refined} : \text{ReprBuf } T \rightarrow \text{ReprBuf } \{t : T \mid P \ t\} \end{array} \right\}$$

Every such type comes with a projection function to the `Buf` type

$$\begin{aligned} \text{as-buf} &: \{r : \text{ReprBuf } T\} \rightarrow T \rightarrow \text{Buf} \\ \text{as-buf } \{r = \text{buf}\} &x = x \\ \text{as-buf } \{r = \text{from } t\} &x = \text{as-buf } t \ (\text{repr } x) \\ \text{as-buf } \{r = \text{refined } t\} &(x, _) = \text{as-buf } t \ x \end{aligned}$$

which eventually computes to the identity function after applying `repr` the appropriate amount of times. Upon compilation, every type is converted to its eventual representation, and all `repr` calls are erased, so the `as-buf` function becomes the identity function at runtime, given that the `r` argument is known at compile-time and monomorphised.

3 A type system for data representations

In this section, we will develop an extension of dependent type theory with inductive families and custom data representations. We start in section 3.2 by exploring the semantics of data representations in terms of algebras for signatures. In section 3.5 we define a core language with inductive families `DATATT` with data representations. The base theory is intensional Martin-Löf type theory (MLTT) [23] with a single universe $\mathcal{U} : \mathcal{U}$. We omit considerations of consistency and universe hierarchy, though these can be added if needed. All of the examples in the paper so far are written in a surface language that elaborates to `DATATT`. In section 3.5, we define a modality `Repr` that allows us to convert between inductive types and their representations. Finally, in section 3.6 we define a translation from `DATATT` to *extensional* MLTT, which ‘elaborates away’ all inductive families to their representations.

The languages we work with are defined in an intrinsically well-formed manner as a setoid over definitional equality, with de-Brujin indices for variables. Weakening of terms is generally left implicit to reduce syntactic noise, and often named notation is used when indices are implied. We use $(a : A) \rightarrow B$ for dependent functions, $(a : A) \times B$ for dependent pairs, $a \equiv_A a'$ for propositional equality, and $a = a' : A$ for definitional equality. Substitution is denoted with

square brackets: if $\Gamma, A \vdash B$ and $\Gamma \vdash a : A$ then $\Gamma \vdash B[a]$. We also notationally identify elements of the universe $A : \mathcal{U}$ with types A `type`, though explicit coercions `El` and `code` between them are implied. Besides the usual judgement forms of MLTT, we also have telescopic judgement forms

$$\begin{aligned} \Gamma \vdash \Delta \text{ tel} \quad & \Delta \text{ is a telescope in } \Gamma, \\ \Gamma \vdash \delta :: \Delta \quad & \delta \text{ is a spine (list of terms) matching telescope } \Delta, \end{aligned}$$

with accompanying rules shown in fig. 1.

$$\begin{array}{c} \text{TEL-EMPTY} \\ \hline \Gamma \vdash \bullet \text{ tel} \end{array} \qquad \begin{array}{c} \text{TEL-EXTEND} \\ \Gamma \vdash A \text{ type} \quad \Gamma, A \vdash \Delta \text{ tel} \\ \hline \Gamma \vdash (A, \Delta) \text{ tel} \end{array}$$

$$\begin{array}{c} \text{SPINE-EMPTY} \\ \hline \Gamma \vdash () :: \bullet \end{array} \qquad \begin{array}{c} \text{SPINE-EXTEND} \\ \Gamma \vdash a : A \quad \Gamma \vdash \delta : \Delta[a] \\ \hline \Gamma \vdash (a, \delta) :: (A, \Delta) \end{array}$$

Fig. 1. Rules for forming telescopes and spines.

Extending contexts by telescopes (such as Γ, Δ) is defined by induction on telescopes. We write $\Delta \rightarrow X$ for an iterated function type with codomain $\Gamma, \Delta \vdash X$, and $(\delta :: \Delta) \rightarrow X[\delta]$ when names are highlighted. We will often use the notation $\delta.y$ to extract a certain index y from a spine δ . This is used when we define telescopes using named notation. For example, if $\delta :: (X : A \rightarrow \mathcal{U}, y : (a : A) \rightarrow X a)$, then $\delta.X : A \rightarrow \mathcal{U}$ and $\delta.y : (a : A) \rightarrow \delta.X a$.

3.1 Algebraic signatures

A representation of a data type must be able to emulate the behaviour of the original data type. In turn, the behaviour of the original data type is determined by its elimination, or induction principle. This means that a representation of a data type should provide an implementation of induction of the same ‘shape’ as the original. Induction can be characterised in terms of algebras and displayed algebras of algebraic signatures [6,22]. Algebraic signatures consist of a list of operations, each with a specified arity. There are many flavours of algebraic signatures with varying degrees of expressiveness. For this paper, we are interested in the ones which can be used as a syntax for defining inductive families in a type theory. Thus, we define two new judgement forms

$$\begin{aligned} \Gamma \vdash S \text{ sig } \Delta \quad & S \text{ is a signature with indices } \Delta \text{ in context } \Gamma, \\ \Gamma \vdash O \text{ op } \Delta \quad & O \text{ is an operation with indices } \Delta \text{ in context } \Gamma. \end{aligned}$$

with accompanying rules shown in fig. 1. Signatures are lists of operations, and operations build up constructor types.

$$\begin{array}{c}
\text{SIG-EMPTY} \\
\frac{\Gamma \vdash \Delta \text{ tel}}{\Gamma \vdash \epsilon \text{ sig } \Delta} \\
\\
\text{SIG-EXTEND} \\
\frac{\Gamma \vdash \Delta \text{ tel} \quad \Gamma \vdash O \text{ op } \Delta \quad \Gamma \vdash S \text{ sig } \Delta}{\Gamma \vdash (O \triangleleft S) \text{ sig } \Delta} \\
\\
\text{OP-EXT} \quad \frac{\Gamma \vdash A \text{ type} \quad \Gamma, A \vdash O \text{ op } \Delta}{\Gamma \vdash (A \rightarrow_{\text{ext}} O) \text{ op } \Delta} \quad
\text{OP-INT} \quad \frac{\Gamma \vdash \delta :: \Delta \quad \Gamma \vdash O \text{ op } \Delta}{\Gamma \vdash (\iota \delta \rightarrow_{\text{int}} O) \text{ op } \Delta} \quad
\text{OP-RET} \quad \frac{\Gamma \vdash \delta :: \Delta}{\Gamma \vdash (\iota \delta) \text{ op } \Delta}
\end{array}$$

Fig. 2. Rules for forming signatures and operations.

Each signature is described by an associated telescope of indices Δ , and a *finite list* of operations:

- $A \rightarrow_{\text{ext}} O$, a (dependent) abstraction over some external type A , of another operation O .
- $\iota \delta \rightarrow_{\text{int}} O$, an abstraction over a recursive occurrence of the object being defined, with indices δ , of another operation O .
- $\iota \delta$, a constructor of the object being defined, with indices δ .

Example 1 (Natural numbers). The signature for natural numbers is indexed by the empty telescope \bullet . It is defined as $\Gamma \vdash (\iota () \triangleleft \iota () \rightarrow_{\text{int}} \iota () \triangleleft \epsilon) \text{ sig } \bullet$. We can add labels to aid readability, omit index spines if they are empty, and omit the final ϵ from signatures:

$$\Gamma \vdash (\text{zero} : \iota \triangleleft \text{succ} : \iota \rightarrow_{\text{int}} \iota) \text{ sig } \bullet.$$

Example 2 (Vectors). The signature for vectors of elements of type T and length n is indexed by the telescope $(T : \mathcal{U}, n : \mathbb{N})$, defined as

$$\begin{aligned}
&\Gamma \vdash (\text{nil} : (T : \mathcal{U}) \rightarrow_{\text{ext}} \iota T \text{ zero} \triangleleft \\
&\quad \text{cons} : (T : \mathcal{U}) \rightarrow_{\text{ext}} (n' : \mathbb{N}) \rightarrow_{\text{ext}} (t : T) \rightarrow_{\text{ext}} \iota T n' \rightarrow_{\text{int}} \iota T (\text{succ } n')) \\
&\quad \text{sig } (T : \mathcal{U}, n : \mathbb{N}).
\end{aligned}$$

Later we will see how we can use the signature in example 1 to define the type of natural numbers $\Gamma \vdash \mathbb{N} \text{ type}$.

Notice that this syntax only allows occurrences of ι in positive positions, which is a requirement for inductive types. Different classes of algebraic signatures, theories and quantification are explored in detail by Kovács [22]. We make no distinction between parameters and indices, though it is possible to add parameters by augmenting the syntax for signatures with an extra telescope that must be uniform across operations.

3.2 Interpreting signatures in the type theory

In order to make use of our definition for algebraic signatures, we would like to be able to interpret their structure as types in the type theory.

Algebras An *algebra* for a signature $\Gamma \vdash S \text{ sig } \Delta$ and carrier type $\Gamma, \Delta \vdash X \text{ type}$ interprets the structure of S in terms of the type X . Concretely, this produces a telescope which matches the structure of S but replaces each occurrence of $\iota \delta$ with $X[\delta]$. The function arrows \rightarrow_{int} and \rightarrow_{ext} in S are interpreted as the function arrow \rightarrow of the type theory.

Example 3 (Natural numbers). An algebra for the signature of natural numbers (example 1) over a carrier $\Gamma \vdash N \text{ type}$ is a spine matching the telescope

$$\Gamma \vdash (\text{zero} : N, \text{succ} : N \rightarrow N) \text{ tel.}$$

Example 4 (Vectors). An algebra for the signature of vectors (example 4) over a carrier $\Gamma, T : \mathcal{U}, n : \mathbb{N} \vdash V \text{ type}$ is a spine matching the telescope

$$\begin{aligned} \Gamma \vdash (\text{nil} : (T : \mathcal{U}) \rightarrow V[T, \text{zero}], \\ \text{cons} : (T : \mathcal{U}) \rightarrow (n' : \mathbb{N}) \rightarrow (t : T) \rightarrow (ts : V[T, n']) \rightarrow V[T, \text{succ } n']) \text{ tel.} \end{aligned}$$

Induction The actual type of natural numbers $\Gamma \vdash \mathbb{N} \text{ type}$ is just the carrier of an algebra over the signature of natural numbers. In particular, the ‘best’ such algebra: one whose operations do not forget any information. In the language of category theory, this is the initial algebra in the category of algebras over the signature of natural numbers. An equivalent formulation of initial algebras is algebras which support *induction*, which is more suitable for our (syntactic) purposes. An algebra $\alpha :: (\text{zero} : X, \text{succ} : X \rightarrow X)$ for natural numbers supports induction if:

For any type family $X \vdash Y \text{ type}$, if we can construct a $\text{zero}_Y : Y[\alpha.\text{zero}]$ and a $\text{succ}_Y : (x : X) \rightarrow Y[x] \rightarrow Y[\alpha.\text{succ } x]$, then we can construct a $\sigma[x] : Y[x]$ for all $x : X$.

The type family Y is commonly called the *motive*, and $(\text{zero}_Y, \text{succ}_Y)$ are the *methods*. The produced term family $x : X \vdash \sigma : Y[x]$ is a *section* of the type family Y . Induction also requires that the section acquires its values from the provided methods. This means that

$$\sigma[\alpha.\text{zero}] = \text{zero}_Y \quad \sigma[\alpha.\text{succ } x] = \text{succ}_Y x \sigma[x].$$

We call these *coherence conditions*. A section that satisfies these conditions is called a *coherent section*. These equations might or might not hold definitionally. In the former case, we have the definitional equalities

$$\begin{aligned} \Gamma \vdash \sigma[\alpha.\text{zero}] &= \text{zero}_Y : Y[\alpha.\text{zero}] \\ \Gamma, x : X \vdash \sigma[\alpha.\text{succ } x] &= \text{succ}_Y x \sigma[x] : Y[\alpha.\text{succ } x]. \end{aligned}$$

In the latter case, we have a spine of propositional equality witnesses

$$\begin{aligned} \Gamma \vdash \sigma_{\text{coh}} :: (\text{zero}_{\text{coh}} : \sigma[\alpha.\text{zero}] \equiv \text{zero}_Y, \\ \text{succ}_{\text{coh}} : (x : X) \rightarrow \sigma[\alpha.\text{succ } x] \equiv \text{succ}_Y x \sigma[x]). \end{aligned}$$

Displayed algebras Notice that the methods $(zero_Y, succ_Y)$ look like an algebra for the signature of natural numbers too, but their carrier is now a type family over another algebra carrier X , and the types of the operations mention both X and Y , using α to go from Δ to X . These are *displayed algebras*. In general, a displayed algebra for a signature $\Gamma \vdash S \text{ sig } \Delta$, algebra α for S over carrier $\Gamma, \Delta \vdash X \text{ type}$, and carrier family $\Gamma, \Delta, X \vdash Y \text{ type}$, interprets the structure of S in terms of both X and Y . This produces a telescope which matches the structure of S but replaces each recursive occurrence $\iota \delta$ with an argument $x : X$ as well as an argument $y : Y[x]$. Each operation returns a Y with indices computed from α . Again, the function arrows in S are interpreted as function types in the type theory.

Example 5 (Vectors). A displayed algebra for an algebra $(nil, cons)$ for vectors (example 4) over a carrier family $\Gamma, T : \mathcal{U}, n : \mathbb{N}, v : V[T, n] \vdash W \text{ type}$ is a spine matching the telescope

$$\begin{aligned} \Gamma \vdash (nil_W : (T : \mathcal{U}) \rightarrow W[T, zero, nil\ T]), \\ cons_W : (T : \mathcal{U}) \rightarrow (n' : \mathbb{N}) \rightarrow (t : T) \rightarrow (ts : V[T, n']) \\ \rightarrow (ts_W : W[T, n', ts]) \rightarrow W[T, succ\ n', cons\ T\ n'\ t\ ts]) \text{ tel.} \end{aligned}$$

In practice, in a call-by-value setting, it is desirable for the inductive hypotheses of a displayed algebra $(ts_W$ above) to be *lazy* values. This improves performance when the inductive hypotheses are not needed. We leave this as an implementation detail.

Finally, we come to the central definition that classifies the algebras which support induction:

Definition 1. *An algebra is inductive if every displayed algebra over it has a coherent section.*

The elimination rule for inductive data types in programming languages is exactly this: given any motive and methods (a displayed algebra), we get a dependent function from the type of the scrutinee to the type of the motive (a section). Furthermore this function satisfies some appropriate computation rules: when we plug in a constructor, we get the result of the method corresponding to it (the coherence conditions). Usually in programming languages, these conditions hold definitionally, as they are the primary means of computation with data.

3.3 Defining algebras and friends

In order to utilise these constructions for our type system, we now explicitly define the following objects:

$\Gamma \vdash S^{\text{alg}} X \text{ tel}$	Algebras for a signature $\Gamma \vdash S \text{ sig } \Delta$ over a carrier $\Gamma, \Delta \vdash X \text{ type}$.
$\Gamma \vdash \alpha^{\text{dispAlg}} Y \text{ tel}$	Displayed algebras for an algebra $\Gamma \vdash \alpha :: S^{\text{alg}} X$ over a motive $\Gamma, \Delta, X \vdash Y \text{ type}$.
$\Gamma \vdash \beta^{\text{coh}} \sigma \text{ tel}$	Propositional coherences for a section $\Gamma, \Delta, X \vdash \sigma : Y$ of a displayed algebra $\Gamma \vdash \beta :: \alpha^{\text{dispAlg}} Y$.

All constructions labelled with superscripts are not part of the syntax of the type system, but rather functions in the metatheory which compute syntactic objects such as telescopes.

The algebras for a signature are defined by case analysis on S :

$$\boxed{\Gamma \vdash S^{\text{alg}} X \text{ tel}} \quad \epsilon^{\text{alg}} X = \bullet \quad (O \triangleleft S')^{\text{alg}} X = ((\nu :: O^{\text{in}} X) \rightarrow X[\nu^{\text{out}}], S'^{\text{alg}} X)$$

An empty signature ϵ produces an empty telescope, while an extended signature $O \triangleleft S'$ produces a telescope extended with a function corresponding to O . This function goes from the inputs of O interpreted in X , to X evaluated at the output indices. The inputs and outputs of each operation O in an algebra are defined by case analysis on O :

$$\boxed{\Gamma \vdash O^{\text{in}} X \text{ tel}} \quad \boxed{\Gamma \vdash \{O\} \nu^{\text{out}} :: \Delta}$$

$$\begin{aligned} (A \rightarrow_{\text{ext}} O')^{\text{in}} X &= (a : A, O'[a]^{\text{in}} X) & \{O = A \rightarrow_{\text{ext}} O'\} (a, \nu')^{\text{out}} &= \nu'^{\text{out}} \\ (\iota \delta \rightarrow_{\text{int}} O')^{\text{in}} X &= (x : X[\delta], O'^{\text{in}} X) & \{O = \iota \delta \rightarrow_{\text{int}} O'\} (x, \nu')^{\text{out}} &= \nu'^{\text{out}} \\ (\iota \delta)^{\text{in}} X &= \bullet & \{O = \iota \delta\} ()^{\text{out}} &= \delta \end{aligned}$$

A similar construction can be performed for displayed algebras over algebras. Displayed algebras are defined by case analysis on S , which is an implicit parameter of $-\text{dispAlg}$.

$$\boxed{\Gamma \vdash \{S\} \alpha^{\text{dispAlg}} Y \text{ tel}} \quad \{S = \epsilon\} ()^{\text{dispAlg}} Y = \bullet$$

$$\{S = O \triangleleft S'\} (\alpha_O, \alpha')^{\text{dispAlg}} Y = ((\mu :: \alpha_O^{\text{dispIn}} Y) \rightarrow Y[\mu^{\text{dispOut}}], \alpha'^{\text{dispAlg}} Y)$$

We omit the definitions of $-\text{dispIn}$ and $-\text{dispOut}$ and refer to the Agda formalisation, but they are similar to the definitions of $-\text{in}$ and $-\text{out}$.

Next we define the coherence conditions for a displayed algebra as a telescope

$$\boxed{\Gamma \vdash \{S\} \{\alpha\} \beta^{\text{coh}} \sigma \text{ tel}}$$

$$\begin{aligned} \{S = \epsilon\} \{\alpha = ()\} ()^{\text{coh}} \sigma &= \bullet \\ \{S = O \triangleleft S'\} \{\alpha = (\alpha_O, \alpha')\} (\beta_O, \beta')^{\text{coh}} \sigma \\ &= ((\nu :: O^{\text{in}} X) \rightarrow \sigma[\alpha_O \nu] \equiv \beta_O (\sigma \$ \nu), \beta'^{\text{coh}} \sigma) \end{aligned}$$

The notation $\sigma \$ \nu$ applies the section σ to the input ν , yielding a displayed input by sampling the section to get the inductive hypotheses.

Now we can define induction for an algebra α as a type

$$\boxed{\Gamma \vdash \{S\} \alpha^{\text{ind}} \text{ type}}$$

$$\{S\} \alpha^{\text{ind}} = (Y : (\delta :: \Delta) \rightarrow X[\delta] \rightarrow \mathcal{U}) \rightarrow (\beta :: \alpha^{\text{dispAlg}} (\delta. x. Y \delta x))$$

$$\rightarrow (\sigma : (\delta :: \Delta) \rightarrow (x : X[\delta]) \rightarrow Y \delta x) \times (\rho :: \beta^{\text{coh}} (\delta. x. \sigma \delta x)).$$

where Y is the motive, β are the methods, and σ is the output section which must satisfy the propositional coherence conditions ρ . Finally, we can package an inductive algebra over a signature as a telescope

$$\boxed{\Gamma \vdash S^{\text{indAlg}} \text{ tel}}$$

$$S^{\text{indAlg}} = (X : \Delta \rightarrow \mathcal{U}, \alpha :: S^{\text{alg}} (\delta. X \delta), \kappa : \alpha^{\text{ind}}),$$

by collecting the carrier X , algebra α and induction κ all together.

3.4 Constructing inductive families

We now extend intensional MLTT with a type for inductive families, which we denote $\text{data}_\Delta S \gamma$. This type defines an inductive family matching a signature S with indices Δ , together with an inductive algebra γ which ‘implements’ the signature S . Notice that this is different to the usual way that inductive families are defined in type theory, for example W-types [5], where all we need to provide is a signature.³ Here, we must also implement the signature and prove the induction principle by providing γ , rather than it being ‘provided’ by the type theory. For example, if our type theory W-types, then we can implement the signature using W-types. In effect, this will later allow us to ‘elaborate away’ inductive definitions to their defined representations. This leads us to the formal definition of a representation:

Definition 2. *A representation of a signature S is an inductive algebra for S .*

In fig. 3 we define the **data** type, and its corresponding introduction, elimination, and computation rules.

³ Notice that the data defining a W-type $(A : \mathcal{U}, B : A \rightarrow \mathcal{U})$ can be viewed as a kind of signature, where A describes the operations and their non-recursive parameters, while B describes the recursive parameters.

$$\begin{array}{c}
\text{DATA-FORM} \quad \frac{\Gamma \vdash S \text{ sig } \Delta \quad \Gamma \vdash \gamma :: S^{\text{indAlg}} \quad \Gamma \vdash \delta :: \Delta}{\Gamma \vdash \text{data}_\Delta S \gamma \delta \text{ type}} \quad \text{DATA-INTRO} \quad \frac{O \in S \quad \Gamma \vdash \nu :: O^{\text{in}} (\text{data}_\Delta S \gamma)}{\Gamma \vdash \text{ctor}_{S.O} \nu : \text{data}_\Delta S \gamma \nu^{\text{out}}} \\
\\
\text{DATA-ELIM} \quad \frac{\Gamma, \delta :: \Delta, \text{data}_\Delta S \gamma \delta \vdash M \text{ type} \quad \Gamma \vdash \beta :: \text{ctor}_S^{\text{dispAlg}} M \quad \Gamma \vdash \delta :: \Delta \quad \Gamma \vdash x : \text{data}_\Delta S \gamma \delta}{\Gamma \vdash \text{elim}_S M \beta \delta x : M[\delta, x]} \\
\\
\text{DATA-COMP} \quad \frac{O \in S \quad \Gamma \vdash \nu :: O^{\text{in}} (\text{data}_\Delta S \gamma) \quad \Gamma, \delta :: \Delta, \text{data}_\Delta S \gamma \delta \vdash M \text{ type} \quad \Gamma \vdash \beta :: \text{ctor}_S^{\text{dispAlg}} M}{\Gamma \vdash \text{elim}_S M \beta \nu^{\text{out}} (\text{ctor}_{S.O} \nu) = \beta_O (\text{elim}_S M \beta \$ \nu) : M[\nu^{\text{out}}, \text{ctors}_{S.O} \nu]}
\end{array}$$

Fig. 3. Rules for data types, constructors and eliminators. We write $O \in S$ to indicate that O is an operation in the signature S . We write α_O to extract the telescope element corresponding to operation O from the algebra α for S .

Constructors form an algebra for the signature S over $\text{data}_\Delta S \gamma$, denoted by $\text{ctor}_S = (\text{ctor}_{S.O})_{O \in S}$. Similarly, the eliminator forms a coherent section over the constructor algebra, which holds definitionally.

One might think, what do we gain by adding **data** to the theory? If we can provide an inductive algebra γ for a signature S ourselves, then why not just use γ directly? The reason is that by having a primitive for inductive types, we can take advantage of their properties in an extensional way. For example, an induction principle suggests that the constructors corresponding to each method are disjoint. Since constructors **ctor** are a primitive term in the theory, we can make use of this when formulating a unification algorithm. Aside from disjointness, we can also rely on other properties such as injectivity, acyclicity, and ‘no-confusion’. McBride [24] originally explored the properties which arise from the existence of induction principles, or equivalently, initiality.

For example, if we have an inductive algebra $(N, \text{zero}_N, \text{succ}_N, \text{elim}_N)$ for the natural numbers signature NatSig , we can prove propositionally that for all $x : N$, $\text{zero}_N \neq \text{succ}_N x$, by invoking elim_N . However, the typechecker does not know this fact; it is not derivable as a definitional equality contradiction. However, it *is* derivable definitionally that for all $x : \mathbb{N}$, $\text{ctor}_{\text{zero}} \neq \text{ctor}_{\text{succ}} x$, where $\mathbb{N} = \text{data NatSig } (N, \text{zero}_N, \text{succ}_N, \text{elim}_N)$ because the metatheoretic syntax does not equate ctor_i and ctor_j unless $i = j$.

Another way to state this is that the existence of **data** enables the use of *dependent pattern matching* on its inhabitants. Nested pattern matching on \mathbb{N} , for example, can be elaborated to invocations of $\text{elim}_\mathbb{N}$, which has the expected computation rules as shown in fig. 3. Converting dependent pattern matching to eliminators has been explored in depth by Goguen, McBride and McKinna [19],

as well as by Cockx and Devriese [16] in absence of Axiom K. These methods can be used to implement pattern matching on `data`.

3.5 Reasoning about representations

So far we are able to construct data types using the $\text{data}_\Delta S \gamma$ type constructor. These data types are themselves implemented in terms of inductive algebras. However, the rules for data types do not utilise them. We would like to be able to relate data types to their underlying inductive algebras. One reason is to avoid unnecessary computation. If we have a type X that is the carrier of two inductive algebras (X, α, κ) and (X, α', κ') for signatures S and S' respectively, then we can form the data types $D = \text{data}_\Delta S (X, \alpha, \kappa)$ and $D' = \text{data}_\Delta S' (X, \alpha', \kappa')$ and make use of the structural properties of initiality. However, we would also like to be able to freely convert between X , D and D' without incurring any runtime cost. After all, D and D' are meant to be translated away to their underlying representation, X . This argument can also be made in the context of theorem proving: sometimes it is easier to prove a property about D or D' , due to their structure, but we should be able to ‘transport’ the property to X .

To make use of these conversions in a computationally-irrelevant manner, while still retaining the fact that D , D' and X are distinct types, we introduce a modality

$$\text{Repr} : \mathcal{U} \rightarrow \mathcal{U},$$

which takes types to their representations. It comes with two term formers `repr` and `unrepr`, which are definitional inverses of each other. We highlight the main rules of `Repr` in fig. 4.

$$\begin{array}{c}
\begin{array}{c} \text{REPR-FORM} \\ \hline \Gamma \vdash A \text{ type} \\ \hline \Gamma \vdash \text{Repr } A \text{ type} \end{array} \qquad
\begin{array}{c} \text{REPR-INTRO} \\ \hline \Gamma \vdash a : A \\ \hline \Gamma \vdash \text{repr } a : \text{Repr } A \end{array} \qquad
\begin{array}{c} \text{REPR-ELIM} \\ \hline \Gamma \vdash a : \text{Repr } A \\ \hline \Gamma \vdash \text{unrepr } a : A \end{array} \\[10pt]
\begin{array}{c} \text{REPR-ID}_1 \\ \hline \Gamma \vdash a : \text{Repr } A \\ \hline \Gamma \vdash \text{repr } (\text{unrepr } a) = a : \text{Repr } A \end{array} \qquad
\begin{array}{c} \text{REPR-ID}_2 \\ \hline \Gamma \vdash a : A \\ \hline \Gamma \vdash \text{unrepr } (\text{repr } a) = a : A \end{array} \\[10pt]
\begin{array}{c} \text{REPR-DATA} \\ \hline \Gamma \vdash S \text{ sig } \Delta \quad \Gamma \vdash \gamma :: S^{\text{indAlg}} \quad \Gamma \vdash \delta :: \Delta \\ \hline \Gamma \vdash \text{Repr } (\text{data}_\Delta S \gamma \delta) = \gamma.X \delta \end{array}
\end{array}$$

Fig. 4. Introduction and elimination forms, as well as computation rules for the `Repr` modality.

These rules allow us to go between a data type $D = \text{data}_\Delta S \gamma \delta$ and its representation $\gamma.X \delta$. In the translation to extensional MLTT that we are yet

to define, this modality is also translated away. Indeed, its purpose is purely intensional: we do not want to equate $\mathbf{data}_\Delta S \gamma$ with $\gamma.X$ because that would render conversion checking undecidable, but we still want to make use of the fact that these types are ‘the same’. In other words, $\mathbf{Repr} A \simeq A$ but not $\mathbf{Repr} A = A$. All the contextual machinery of general modal type systems [20] is not necessary here because this modality is fibred over contexts so it presents as just a type former.

One might hope for additional computation rules. For example, the representation of a constructor should be equal to the underlying algebra element of the constructor type’s representation:⁴

$$\mathbf{repr} (\mathbf{ctor}_{S.O} \nu) = \gamma.\alpha_O (\mathbf{repr} \nu).$$

Unfortunately, having this as a computation rule would render conversion checking undecidable, because if one applies \mathbf{unrepr} to a term $\mathbf{repr} (\mathbf{ctor}_{S.O} \nu)$ which has already been reduced to its representation, $\mathbf{unrepr} (\gamma.\alpha_O (\mathbf{repr} \nu))$, there is no clear way to decide that this is convertible to $\mathbf{ctor}_{S.O} \nu$ even though the definitional equality rules would imply that it is (due to $\mathbf{REPR-ID}_2$). Nevertheless, we can still postulate this equality propositionally, and it is justified by the translation step.

We should also ask for compatibility equations between \mathbf{Repr} and the rest of the type formers of MLTT; for example $\mathbf{Repr} ((a : A) \times B[a]) = (a : \mathbf{Repr} A) \times \mathbf{Repr} B[\mathbf{unrepr} a]$, which can hold definitionally without breaking conversion. These are given for \mathcal{U} , Π , Σ , unit and identity types in the Agda formalisation.

Subuniverse of concrete types As described, the modality \mathbf{Repr} is not definitionally idempotent: $\mathbf{Repr} (\mathbf{Repr} A) = \mathbf{Repr} A$ does not always hold. An inductive algebra used as a representation of a data type might itself be implemented in terms of another data type, which again reduces under the action of \mathbf{Repr} . A more principled approach might be to view the image of \mathbf{Repr} as a subuniverse of \mathcal{U} . The restriction

$$\mathbf{Repr} : \mathcal{U} \rightarrow \mathcal{U}_C$$

targets a universe of ‘concrete’ types $\mathcal{U}_C < \mathcal{U}$ closed under all standard type formers, but which does not contain any \mathbf{data} types. We can then require that all inductive algebras γ used in the rule $\mathbf{DATA-FORM}$ must have a concrete carrier $X : \Delta \rightarrow \mathcal{U}_C$. This does not limit expressivity because we can always wrap any inductive algebra carrier with \mathbf{Repr} to bring it down to \mathcal{U}_C . We do not assume this additional structure for simplicity, but it might be a useful feature in practice. For example, it would simplify the transitivity example in section 2.5.

3.6 Translating to extensional MLTT

We now define a type- and equality-preserving translation step \mathcal{R} shown in fig. 5 from DATATT to extensional MLTT, meant to be applied during the compilation

⁴ When we write $\mathbf{repr} \nu$ we mean to apply \mathbf{repr} to all recursive occurrences (all places that ι appears in the domain of O).

process. The extensional flavour of MLTT involves adding the equality reflection rule

$$\frac{\text{REFLECT } p \quad \Gamma \vdash p : a =_A a'}{\Gamma \vdash a \equiv a' : A}$$

General undecidability of conversion is not a problem because type checking is decidable for DATATT⁵ and we only need to apply this transformation after type checking, on fully-typed terms. The translation is defined over the syntax of DATATT [13] such that definitional equality is preserved, which is shown in fig. 5. \mathcal{R} replaces data types with their underlying inductive algebras. We use the notation Ty and Tm for types and terms respectively, with a subscript indicating the language.

$\mathcal{R} : \text{Ty}_{\text{DATATT}} \Gamma \rightarrow \text{Ty}_{\text{MLTT}} \mathcal{R}\Gamma$	$\mathcal{R} : \text{Tm}_{\text{DATATT}} \Gamma A \rightarrow \text{Tm}_{\text{MLTT}} \mathcal{R}\Gamma \mathcal{R}A$
$\mathcal{R}(\text{Repr } A) = \mathcal{R}A$ $\mathcal{R}(\text{data}_{\Delta} S \gamma \delta) = \mathcal{R}(\gamma.X) \mathcal{R}\delta$ (otherwise recurse with \mathcal{R})	$\mathcal{R}(\text{repr } t) = \mathcal{R}t$ $\mathcal{R}(\text{unrepr } t) = \mathcal{R}t$ $\mathcal{R}(\text{ctor}_{S.O} \{\gamma\} \nu) = \mathcal{R}(\gamma.\alpha_O) \mathcal{R}\nu$ $\mathcal{R}(\text{elim}_S \{\gamma\} M \beta \delta x)$ $\quad = (\mathcal{R}(\gamma.\kappa) \mathcal{R}M \mathcal{R}\beta).\sigma \mathcal{R}\delta \mathcal{R}x$ (otherwise recurse with \mathcal{R})

Fig. 5. Translation of DATATT to MLTT. It replaces data types with their underlying inductive algebras, and eliminators by the induction principle provided by representations.

All the mappings above are clearly structurally recursive. The translation is extended to contexts, substitutions, telescopes and spines pointwise, and all the rest of the syntax is preserved: $\mathcal{R}((a : A) \rightarrow B[a]) = (a : \mathcal{R}A) \rightarrow (\mathcal{R}B)[a]$ etc. All **data** types are translated to the carriers of their inductive algebras. Invocations of **Repr** are removed. One can view **Repr** as locally applying \mathcal{R} to a part of the program. The definitional equality preservation by \mathcal{R} is shown in fig. 6. The isomorphism of **repr/unrepr**, as well as the rule DATA-REPR are preserved by metatheoretic reflexivity on the other side, since all representation operators are erased. Coherence rules for eliminators are preserved by reflecting the propositional coherence rules provided by their defined representations. We have formalised this transformation in Agda.

⁵ Not formalised in this paper.

$\mathcal{R}_{\approx}^{\text{Ty } \Gamma} : A \approx_{\text{DATATT}} A' \rightarrow \mathcal{R}A \approx_{\text{MLTT}} \mathcal{R}A'$	$\mathcal{R}_{\approx}^{\text{Trm } \Gamma \ A} : a \approx_{\text{DATATT}} a' \rightarrow \mathcal{R}a \approx_{\text{MLTT}} \mathcal{R}a'$
$\mathcal{R}(\text{REPR-DATA } \{S, \gamma, \Delta\})$ $\quad = \text{REFL } (\mathcal{R}(\gamma.X) \ \mathcal{R}\delta)$ (otherwise recurse with \mathcal{R}_{\approx})	$\mathcal{R}(\text{REPR-ID}_1 \ \{a\}) = \text{REFL } \mathcal{R}a$ $\mathcal{R}(\text{REPR-ID}_2 \ \{a\}) = \text{REFL } \mathcal{R}a$ $\mathcal{R}(\text{DATA-COMP } \{S, O, \gamma, M, \beta, \nu\})$ $\quad = \text{REFLECT } (\mathcal{R}(\gamma.\kappa) \ \mathcal{R}M \ \mathcal{R}\beta). \rho_O \ \mathcal{R}\nu$ (otherwise recurse with \mathcal{R}_{\approx})

Fig. 6. Equality translation of DATATT to MLTT. This amounts to an inductive proof that \mathcal{R} preserves equality. The computation rule for eliminators is translated by reflecting the coherence of sections produced by applying the induction on representations.

Theorem 1. \mathcal{R} preserves typing and definitional equality.

Proof. By the setoid homomorphism $(\mathcal{R}, \mathcal{R}_{\approx})$.

\mathcal{R} is not injective in general, because two data types with distinct signatures might have the same representation.

Theorem 2. \mathcal{R} is a left-inverse of the evident inclusion $i : \text{MLTT} \hookrightarrow \text{DATATT}$:

$$\frac{\Gamma \vdash a : A}{\Gamma \vdash \mathcal{R}(i(a)) = a : A} \quad (\text{in MLTT}).$$

Proof. The inclusion produces types without **data** or **Repr**. Thus, the action of \mathcal{R} on the image of i does not invoke the equality reflection rule. With that constraint, and by induction on the syntax, $\mathcal{R} \circ i$ is the identity function on MLTT.

Inductive types in an empty context Basic MLTT with dependent functions, pairs and propositional equality as we have presented here is not sufficient to construct most inductive algebras in an empty context. Without W-types or fixpoints, we must postulate the induction principles we have access to, like the example with GMP integers in section 2. In practice, we very often want to be able to construct inductive types ‘from scratch’. For this, we can simply extend the base theory with W-types or something similar. One especially convenient choice is to extend the syntax of both the source and target languages of \mathcal{R} with a class of data types like in fig. 3, but without requiring them to be implemented by inductive algebras. The advantage is now that we can opt-in when we want to represent inductive types specially, but otherwise fall back to some kind of default implementation. The downside is that now the target still contains inductive types, though this is okay for compilation purposes if we choose a ‘sane default’, usually tagged unions containing indirections. This is the approach we take in SUPERFLUID.

3.7 Computational irrelevance

In order to reason about computational irrelevance, we assume that there is an additional program extraction step from the target of \mathcal{R} into some simply-typed or untyped calculus PROG , denoted by vertical bars $|x|$. As opposed to \mathcal{R} , it might not preserve the definitional equality of the syntax; we might want to compile two definitionally equal terms differently. For example, we might not always want to reduce function application redexes. We will use the `monospace` font for terms in PROG .

Definition 3. A function $\Gamma \vdash f : (a : A) \rightarrow B$, is computationally irrelevant if $|\mathcal{R}A| = |\mathcal{R}B|$ and $|\mathcal{R}f| = \text{id}$.

Theorem 3. The type former *Repr* is injective up to internal isomorphism, i.e.

$$\frac{\Gamma \vdash p : \text{Repr } T =_{\mathcal{U}} \text{Repr } T'}{\Gamma \vdash \text{conv}_p : T \simeq T'}, \quad (2)$$

Moreover, this isomorphism is computationally irrelevant.

Proof. For the input proof p , the forward direction conv_p is $\lambda x. \text{unrepr}_{T'} (J \text{id } (\text{repr } x) p)$ and the backward direction conv_p^{-1} is $\lambda x. \text{unrepr}_T (J \text{id } (\text{repr } x) (\text{sym } p))$. The coherences hold definitionally. Applying \mathcal{R} , we get $\lambda x. J \text{id } x p$, which is the identity function after $|-|$.

Zero-cost refinements Consider extending our languages with usage-aware ‘subset’ dependent pairs

$$\frac{\Gamma \vdash A \text{ type} \quad \Gamma, A \vdash B \text{ type}}{\Gamma \vdash \{A \mid B\} \text{ type}}$$

in such a way that *Repr* and \mathcal{R} preserve them, but the extraction step erases the right component, i.e. $|\{A \mid B\}| = |A|$, $|(x, y)| = |x|$ and $|\pi_1 x| = |x|$.⁶ Suppose we have an inductive family $G = \text{data}_I S_G \gamma_G$ over some index type I , and an inductive type $F = \text{data } S_F \gamma_F$ such that G is represented by a refinement $r : F \rightarrow I$,

$$\gamma_G = (\lambda i. \{f : F \mid r f \equiv_I i\}, \alpha, \kappa).$$

Then, we can construct computationally irrelevant functions

$$\begin{aligned} \text{forget}_i : G \ i \rightarrow F & & \text{remember} : (f : F) \rightarrow G \ (r \ f) \\ \text{forget}_i = \lambda g. \pi_1 \ (\text{repr } g) & & \text{remember} = \lambda x. \text{unrepr } (x, \text{refl}). \end{aligned}$$

By reasoning similar to theorem 3, $|\mathcal{R} \text{forget}_i| = |\mathcal{R} \text{remember}| = \text{id}$.

⁶ This can be implemented using quantitative type theory for example.

4 Implementation

SUPERFLUID is a programming language with dependent types with quantities, inductive families and data representations. Its compiler is written in Haskell and the compilation target is JavaScript. After prior to code generation, the \mathcal{R} transformation is applied to the elaborated core program, which erases all inductive constructs with defined representations. Then, a JavaScript program is extracted, erasing all irrelevant data by usage analysis similarly to Idris 2. As a result, with appropriate postulates in the prelude, we are able to represent `Nat` as JavaScript’s `BigInt`, and `List T/SnocList T/Vec T n` as JavaScript’s arrays with the appropriate index refinement, such that we can convert between them without any runtime overhead. The syntax of SUPERFLUID very closely mirrors the syntax given in the first half of this paper. It supports global definitions, inductive families, as well as postulates. Users are able to define custom representations for data types using `repr – as –` blocks.

Currently we do not require proofs of eliminator coherence, but they are straightforward to add. We also treat the rule for representing constructors ($\text{repr } (\text{ctor}_{S.O} \nu) = \gamma.\alpha_O (\text{repr } \nu)$) as definitional in the implementation, at the cost of breaking decidability of equality, but with the benefit of fewer manual transports. We are currently working on adding dependent pattern matching that is elaborated to internal eliminators, so that we can take advantage of the structural unification rules for data types [24]. We have written some of the examples in this paper in SUPERFLUID, which can be found in the `examples` directory. Overall the implementation is a proof of concept, but we expect that our framework can be implemented in an existing language.

5 Related work

Using inductive types as a form of abstraction was first explored by Wadler [27] through *views*. The extension to dependent types was developed by McBride and McKinna [25], as part of the Epigram project. Our system differs from views in the computational content of the abstraction; even with deforestation [28] views are not always zero-cost, but representations are. Atkey [11] shows how to generically derive inductive types which are refinements of other inductive types. This work could be integrated in our system to automatically generate representations for refined data types. Zero-cost data reuse when it comes to refinements of inductive types has been explored in the context of Church encoding in Cedille [18], but does not extend to custom representations.

Work by Allais [7,8] uses a combination of views, erasure by quantitative type theory, and universes of flattened data types to achieve performance improvements when working with serialised data in Idris 2. Our approach differs because we have access to ‘native’ data representations, so we do not need to rely on encodings. Additionally, they rely on heuristic compiler optimisations to erase their views. On the topic of memory layout optimisation, Baudon [12] develops Ribbit, a DSL for the specification of the memory representation of algebraic

data types, which can specify techniques like struct packing and bit-stealing. To our knowledge however, this does not provide control over the indirection introduced by *inductive* types.

Dependently typed languages with extraction features, notably Coq [3] and Agda [1], have some overlapping capabilities with our approach, but they do not provide any of the correctness guarantees. Optimisation tricks such as the Nat-hack, and its generalisation to other types, can emulate a part of our system but are unverified and special casing in the compiler. Since the extended abstract version of this paper, an optimisation was merged into Idris 2 [2] to erase the forgetful and recomputation functions for reindexing list/maybe/number-like types. There also seems to be a demand for this kind of optimisation in Agda [4].

6 Future work

In the future we aim to expand the class of theories we consider, to include quotient-induction, induction-induction and induction-recursion. Representations for quotient-inductive types in particular could give rise to ergonomic ways of computing with more ‘traditional’ data structures such as hash maps or binary search trees. We could program inductively over these structures but extract programs without redundancy in data representation.

There are elements of our formalisation which should be developed further. We did not formulate decidability of equality and normalisation for DATATT, which is needed for typechecking. We have implemented a normalisation-by-evaluation [9] algorithm used in the implementation of SUPERFLUID, but have only sketched that it has the desired properties.

We would like to further explore the relationship of the *Repr* modality with general systems for defining modalities, such as *multi-modal type theory* by Gratzer [20]. Additionally, we expect that metaprogramming with representations is most naturally done in the context of *two-level type theory* [21]. We would like to explore the embedding of DATATT in a two-level type theory, where signatures become types in the meta-fragment. Then we could develop various methods of generating representations internally, for example through algebraic ornaments, without needing to laboriously prove induction principles by hand.

7 Conclusion

This paper addresses some of the inefficiencies of inductive families in dependently typed languages by introducing custom runtime representations that preserve logical guarantees and simplicity of the surface language while optimising performance and usability. These representations are formalised as inductive algebras, and come with a framework for reasoning about them: provably zero-cost conversions between original and represented data. The compilation process guarantees erasure of abstraction layers, translating high-level constructs to their defined implementations. Our hope is that by decoupling logical structure

from runtime representation, ergonomic type-driven correctness can be leveraged without sacrificing performance.

Acknowledgements We thank the anonymous reviewers for their feedback that significantly improved the quality of this paper, Guillaume Allais for helpful comments on the extended abstract, as well as Nathan Corbyn, Ellis Kesterton, and Matthew Pickering for interesting discussions on algebraic theories, ergonomics of inductive types, and pattern synonyms.

References

1. agda2hs documentation — agda2hs documentation. <https://agda.github.io/agda2hs/>, accessed: 2025-2-19
2. Make ‘cons’, ‘nil’, ‘just’ and ‘nothing’ constructors have uniform names by Z-snails · pull request #3486 · idris-lang/Idris2. <https://github.com/idris-lang/Idris2/pull/3486>, accessed: 2025-2-19
3. Program extraction — coq 8.13.2 documentation. <https://coq.inria.fr/doc/v8.13/refman/addendum/extraction.html>, accessed: 2025-2-19
4. Should agda optimise away the erasure from vec to list? · issue #7701 · agda/agda. <https://github.com/idris-lang/Idris2/pull/3486>, accessed: 2025-2-19
5. Abbott, M., Altenkirch, T., Ghani, N.: Representing nested inductive types using W-types. In: Automata, Languages and Programming, pp. 59–71. Lecture notes in computer science, Springer Berlin Heidelberg, Berlin, Heidelberg (2004), https://link.springer.com/chapter/10.1007/978-3-540-27836-8_8
6. Adamek, J., Rosicky, J., Vitale, E.M.: Cambridge tracts in mathematics: Algebraic theories: A categorical introduction to general algebra series number 184. Cambridge University Press, Cambridge, England (18 Nov 2010), <https://www.cambridge.org/academic/subjects/mathematics/logic-categories-and-sets/algebraic-theories-categorical-introduction-general-algebra?format=HB&isbn=9780521119221>
7. Allais, G.: Builtin types viewed as inductive families. In: Programming Languages and Systems. pp. 113–139. Springer Nature Switzerland (2023), http://dx.doi.org/10.1007/978-3-031-30044-8_5
8. Allais, G.: Seamless, correct, and generic programming over serialised data. arXiv [cs.PL] (20 Oct 2023), <http://arxiv.org/abs/2310.13441>
9. Altenkirch, T., Geniet, C.: Big step normalisation for type theory (24 Sep 2020), <https://drops.dagstuhl.de/storage/00lipics/lipics-vol175-types2019/LIPIcs.TYPES.2019.4/LIPIcs.TYPES.2019.4.pdf>
10. Atkey, R.: Syntax and semantics of quantitative type theory. In: Proceedings of the 33rd Annual ACM/IEEE Symposium on Logic in Computer Science. pp. 56–65. LICS ’18, Association for Computing Machinery, New York, NY, USA (9 Jul 2018), <https://doi.org/10.1145/3209108.3209189>
11. Atkey, R., Johann, P., Ghani, N.: When is a type refinement an inductive type? In: Foundations of Software Science and Computational Structures, pp. 72–87. Lecture notes in computer science, Springer Berlin Heidelberg, Berlin, Heidelberg (2011), <https://bentnib.org/inductive-refinement.pdf>

12. Baudon, T., Radanne, G., Gonnord, L.: Bit-stealing made legal: Compilation for custom memory representations of algebraic data types. *Proc. ACM Program. Lang.* **7**(ICFP), 813–846 (31 Aug 2023), <https://doi.org/10.1145/3607858>
13. Boulier, S., Pédro, P.M., Tabareau, N.: The next 700 syntactical models of type theory. In: *Proceedings of the 6th ACM SIGPLAN Conference on Certified Programs and Proofs*. pp. 182–194. CPP 2017, Association for Computing Machinery, New York, NY, USA (16 Jan 2017), <https://doi.org/10.1145/3018610.3018620>
14. Brady, E., McBride, C., McKinna, J.: Inductive families need not store their indices. In: *Types for Proofs and Programs*. pp. 115–129. Springer Berlin Heidelberg (2004), http://dx.doi.org/10.1007/978-3-540-24849-1_8
15. Castellan, S., Clairambault, P., Dybjer, P.: Categories with families: Untyped, simply typed, and dependently typed. *arXiv [cs.LO]* (1 Apr 2019), <http://arxiv.org/abs/1904.00827>
16. Cockx, J., Devriese, D.: Proof-relevant unification: Dependent pattern matching with only the axioms of your type theory. *J. Funct. Prog.* **28**(e12), e12 (Jan 2018), <https://www.cambridge.org/core/services/aop-cambridge-core/content/view/E54D56DC3F5D5361CCDECA824030C38E/S095679681800014Xa.pdf/div-class-title-proof-relevant-unification-dependent-pattern-matching-with-only-the-axioms-of-your-type-theory.pdf>
17. Dagand, P.E., McBride, C.: A categorical treatment of ornaments. *arXiv [cs.PL]* (16 Dec 2012), <http://arxiv.org/abs/1212.3806>
18. Diehl, L., Firsov, D., Stump, A.: Generic zero-cost reuse for dependent types. *Proc. ACM Program. Lang.* **2**(ICFP), 1–30 (30 Jul 2018), <https://doi.org/10.1145/3236799>
19. Goguen, H., McBride, C., McKinna, J.: Eliminating dependent pattern matching. In: *Algebra, Meaning, and Computation*, pp. 521–540. Lecture notes in computer science, Springer Berlin Heidelberg, Berlin, Heidelberg (2006), <https://research.google.com/pubs/archive/99.pdf>
20. Gratzner, D., Kavvos, G.A., Nuyts, A., Birkedal, L.: Multimodal dependent type theory. *arXiv [cs.LO]* (30 Nov 2020), [http://dx.doi.org/10.46298/lmcs-17\(3:11\)2021](http://dx.doi.org/10.46298/lmcs-17(3:11)2021)
21. Kovács, A.: Staged compilation with two-level type theory. *arXiv [cs.PL]* (20 Sep 2022), <http://arxiv.org/abs/2209.09729>
22. Kovács, A.: Type-theoretic signatures for algebraic theories and inductive types. Ph.D. thesis (2023), https://andraskovacs.github.io/pdfs/phdthesis_compact.pdf
23. Martin-Löf, P.: Intuitionistic type theory **1**, 1–91 (1984), <http://people.csail.mit.edu/jgross/personal-website/papers/academic-papers-local/Martin-Lof80.pdf>
24. McBride, C., Goguen, H., McKinna, J.: A few constructions on constructors. In: *Lecture Notes in Computer Science*, pp. 186–200. Lecture notes in computer science, Springer Berlin Heidelberg, Berlin, Heidelberg (2006), <http://www.e-pig.org/downloads/concon.pdf>
25. McBride, C., McKinna, J.: The view from the left. *J. Funct. Programming* **14**(1), 69–111 (Jan 2004), <https://www.cambridge.org/core/services/aop-cambridge-core/content/view/F8A44CAC27CCA178AF69DD84BC585A2D/S0956796803004829a.pdf/div-class-title-the-view-from-the-left-div.pdf>
26. Moon, B., Eades, III, H., Orchard, D.: Graded modal dependent type theory. In: *Programming Languages and Systems*. pp. 462–490. Springer International Publishing (2021), http://dx.doi.org/10.1007/978-3-030-72019-3_17

27. Wadler, P.: Views: a way for pattern matching to cohabit with data abstraction. In: Proceedings of the 14th ACM SIGACT-SIGPLAN symposium on Principles of programming languages. pp. 307–313. POPL '87, Association for Computing Machinery, New York, NY, USA (1 Oct 1987), <https://doi.org/10.1145/41625.41653>
28. Wadler, P.: Deforestation: transforming programs to eliminate trees. Theor. Comput. Sci. **73**(2), 231–248 (1 Jun 1990), <https://www.sciencedirect.com/science/article/pii/030439759090147A>

8 Appendix

8.1 Implementation

The implementation of SUPERFLUID can be found at <https://github.com/kontheocharis/superfluid>.

8.2 Basic MLTT

We define the basic formulation of Martin-Löf Type Theory with $\mathcal{U} : \mathcal{U}$, Π , Σ , propositional equality and unit. We omit the definition of the substitution calculus and equality coercions (see [15, 5.1.2]). We use de-Brujin indices, keeping weakening implicit, and abuse notation for substitutions of terms: $A[t]$ for $A[\text{id}, t]$.

$\frac{}{\Gamma \vdash \mathcal{U} \text{ type}} \text{UNIV-FORM}$	$\frac{\Gamma \vdash a : \mathcal{U}}{\Gamma \vdash \text{El } a \text{ type}} \text{EL-FORM}$	$\frac{}{\Gamma \vdash \text{code } A : \mathcal{U}} \text{UNIV-INTRO}$	$\frac{\Gamma \vdash A \text{ type} \quad \Gamma, A \vdash B \text{ type}}{\Gamma \vdash A \rightarrow B \text{ type}} \text{PI-FORM}$
	$\frac{\Gamma, A \vdash b : B}{\Gamma \vdash \lambda b : A \rightarrow B} \text{PI-INTRO}$		$\frac{\Gamma \vdash f : A \rightarrow B \quad \Gamma \vdash a : A}{\Gamma \vdash f a : B[a]} \text{PI-ELIM}$
	$\frac{\Gamma \vdash A \text{ type} \quad \Gamma \vdash a : A \quad \Gamma \vdash b : A}{\Gamma \vdash a \equiv_A b \text{ type}} \text{EQ-FORM}$		$\frac{\Gamma \vdash a : A}{\Gamma \vdash \text{refl } a : a \equiv_A a} \text{EQ-INTRO}$
			$\frac{\Gamma \vdash A \text{ type} \quad \Gamma, a : A, b : A, a \equiv_A b \vdash P \text{ type} \quad \Gamma, a : A \vdash r : P[a, a, \text{refl } a] \quad \Gamma \vdash a : A \quad \Gamma \vdash b : A \quad \Gamma \vdash p : a \equiv_A b}{\Gamma \vdash \text{J } P d p : P[a, b, p]} \text{EQ-ELIM}$
$\frac{}{\Gamma \vdash \top \text{ type}} \text{UNIT-FORM}$	$\frac{}{\Gamma \vdash \text{tt} : \top} \text{UNIT-INTRO}$		$\frac{\Gamma \vdash A \text{ type} \quad \Gamma, A \vdash B \text{ type}}{\Gamma \vdash A \times B \text{ type}} \text{SIGMA-FORM}$
$\frac{\Gamma \vdash a : A \quad \Gamma \vdash b : B[a]}{\Gamma \vdash (a, b) : A \times B} \text{SIGMA-INTRO}$		$\frac{\Gamma \vdash p : A \times B}{\Gamma \vdash \text{fst } p : A} \text{SIGMA-ELIM-FST}$	$\frac{\Gamma \vdash p : A \times B}{\Gamma \vdash \text{snd } p : B[\text{fst } p]} \text{SIGMA-ELIM-SND}$

Fig. 7. Typing rules for MLTT.

$\text{El}(\text{code } A) = A$	$x = \text{tt}$
$\text{code}(\text{El } t) = t$	$\text{fst}(a, b) = a$
$(\lambda b) a = b[a]$	$\text{snd}(a, b) = b$
$\lambda x. (f x) = f$	$(\text{fst } p, \text{snd } p) = p$
$\text{J } P r (\text{refl } a) = r a$	

Fig. 8. Homogenous equality rules for MLTT, omitting substitution rules such as $(\text{El } a)[\sigma] = \text{El}(a[\sigma])$.

8.3 Extension to DATATT

The language DATATT is the extension of MLTT as presented above, by the rules in figs. 3 and 4. All constructions of signatures and operations are parametric over contexts, and thus stable under substitution. Below we present the definitional equality rules of Repr.

$\frac{\text{REPR-R}}{\Gamma \vdash \text{unrepr} (\text{repr } t) = t}$	$\frac{\text{REPR-L}}{\Gamma \vdash \text{repr} (\text{unrepr } t) = t}$
$\frac{\text{REPR-PI}}{\Gamma \vdash \text{Repr} (\Pi T U) = \Pi T (\text{Repr } U)}$	$\frac{\text{REPR-LAM}}{\Gamma \vdash \text{repr} (\lambda u) = \lambda (\text{repr } u)}$
$\frac{\text{UNREPR-LAM}}{\Gamma \vdash \text{unrepr} (\lambda u) = \lambda (\text{unrepr } u)}$	$\frac{\text{REPR-APP}}{\Gamma \vdash \text{repr} (f @) = (\text{repr } f) @}$
$\frac{\text{UNREPR-APP}}{\Gamma \vdash \text{unrepr} (f @) = (\text{unrepr } f) @}$	$\frac{\text{REPR-SUB}}{\Gamma \vdash \text{repr} (t[\sigma]) = (\text{repr } t)[\sigma]}$
$\frac{\text{UNREPR-SUB}}{\Gamma \vdash \text{unrepr} (t[\sigma]) = (\text{unrepr } t)[\sigma]}$	$\frac{\text{REPR-TYPE-SUB}}{\Gamma \vdash \text{Repr} (T[\sigma]) = (\text{Repr } T)[\sigma]}$
$\frac{\text{REPR-UNIV}}{\Gamma \vdash \text{Repr } \mathcal{U} = \mathcal{U}}$	$\frac{\text{REPR-CODE}}{\Gamma \vdash \text{repr} (\text{code } T) = \text{code } T}$
$\frac{\text{UNREPR-CODE}}{\Gamma \vdash \text{unrepr} (\text{code } T) = \text{code } T}$	$\frac{\text{REPR-ID}}{\Gamma \vdash \text{Repr} (\text{Id } a b) = \text{Id} (\text{repr } a) (\text{repr } b)}$
$\frac{\text{REPR-REFL}}{\Gamma \vdash \text{repr} (\text{refl } u) = \text{refl} (\text{repr } u)}$	$\frac{\text{UNREPR-REFL}}{\Gamma \vdash \text{unrepr} (\text{refl } u) = \text{refl} (\text{unrepr } u)}$
$\frac{\text{REPR-J}}{\Gamma \vdash \text{repr} (\text{J } C w e) = \text{J} (\text{Repr } C) (\text{repr } w) e}$	
$\frac{\text{UNREPR-J}}{\Gamma \vdash \text{unrepr} (\text{J} (\text{Repr } C) w e) = \text{J } C (\text{unrepr } w) e}$	

8.4 Full translation from MLTT to DATATT

8.5 Displayed algebras over algebras

8.6 Applying sections to algebra inputs

TODO

8.7 Utilities when working with algebras

Realisation functions for algebras